Лабораторная работа №2
по дисциплине
«Методы машинного обучения»
на тему
«Изучение библиотек обработки данных»

Выполнил:
студент группы ИУ5-21М
Красавцев К.Ю.

_____

Москва — 2020 г.

Assignment #1 (demo)

Exploratory data analysis with Pandas

In this task you should use Pandas to answer a few questions about the Adult dataset.

Unique values of all features (for more information, please see the links above):

- `age`: continuous.
- `workclass`: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.
- `fnlwgt`: continuous.
- `education`: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.
- `education-num`: continuous.
- `marital-status`: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.
- `occupation`: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.
- `relationship`: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.
- `race`: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.
- `sex`: Female, Male.
- `capital-gain`: continuous.
- `capital-loss`: continuous.
- `hours-per-week`: continuous.
- `native-country`: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinadad&Tobago, Peru, Hong, Holand-Netherlands.

salary: >50K, <=50K.

3.1 Часть 1

Разведочный анализ данных с Pandas

In [41]:

```python
import pandas as pd
import numpy as np
pd.set_option('display.max.columns', 100)
%matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns
# we don't like warnings
# you can comment the following 2 lines if you'd like to
import warnings
warnings.filterwarnings('ignore')
```

In [43]:

```python
data = pd.read_csv('adult.data.csv')
data.head()
```

Out[43]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black

In [11]:

```python
#1. How many men and women (sex feature) are represented in this dataset?
data['sex'].value_counts()
```

Out[11]:

```
Male      21790
Female    10771
Name: sex, dtype: int64
```

In [12]:

```
#2. What is the average age (age feature) of women?
female_data = data[data['sex'] == 'Female']
```

In [17]:

```
female_data['age'].mean()
```

Out[17]:

36.85823043357163

In [18]:

```
#alt
data.loc[data['sex'] == 'Female', 'age'].mean()
```

Out[18]:

36.85823043357163

In [21]:

```
#3. What is the proportion of German citizens (native-country feature)?
float((data['native-country'] == 'Germany').sum()) / data.shape[0]
```

Out[21]:

0.004207487485028101

In [26]:

```
#4-5. What are mean value and standard deviation of the age
#of those who recieve more than 50K per year (salary feature) and those who receive less than 50K per year?
ages1 = data.loc[data['salary'] == '>50K', 'age']
ages2 = data.loc[data['salary'] == '<=50K', 'age']

print("The average age of the rich: {0} +- {1} years, poor - {2} +- {3} years.".format(
    round(ages1.mean()), round(ages1.std(), 1),
    round(ages2.mean()), round(ages2.std(), 1)))
```

The average age of the rich: 44.0 +- 10.5 years, poor - 37.0 +- 14.0 years.

In [27]:

```
#6. Is it true that people who receive more than 50k have at least high school education?
#(education - Bachelors, Prof-school, Assoc-acdm, Assoc-voc, Masters or Doctorate feature)
data.loc[data['salary'] == '>50K', 'education'].unique()
```

Out[27]:

```
array(['HS-grad', 'Masters', 'Bachelors', 'Some-college', 'Assoc-voc',
       'Doctorate', 'Prof-school', 'Assoc-acdm', '7th-8th', '12th',
       '10th', '11th', '9th', '5th-6th', '1st-4th'], dtype=object)
```

In [31]:

```python
#7. Display statistics of age for each race (race feature) and each gender.
#Use groupby() and describe(). Find the maximum age of men of Amer-Indian-Eskimo race.
#data.groupby(['race', 'sex']).describe()
for (race, sex), sub_df in data.groupby(['race', 'sex']):
    print('Race: {0}, sex {1}'.format(race, sex))
    print(sub_df['age'].describe())
```

```
Race: Amer-Indian-Eskimo, sex Female
count    119.000000
mean      37.117647
std       13.114991
min       17.000000
25%       27.000000
50%       36.000000
75%       46.000000
max       80.000000
Name: age, dtype: float64
Race: Amer-Indian-Eskimo, sex Male
count    192.000000
mean      37.208333
std       12.049563
min       17.000000
25%       28.000000
50%       35.000000
75%       45.000000
max       82.000000
Name: age, dtype: float64
Race: Asian-Pac-Islander, sex Female
count    346.000000
mean      35.089595
std       12.300845
min       17.000000
25%       25.000000
50%       33.000000
75%       43.750000
max       75.000000
Name: age, dtype: float64
Race: Asian-Pac-Islander, sex Male
count    693.000000
mean      39.073593
std       12.883944
min       18.000000
25%       29.000000
50%       37.000000
75%       46.000000
max       90.000000
Name: age, dtype: float64
Race: Black, sex Female
count    1555.000000
mean       37.854019
std        12.637197
min        17.000000
25%        28.000000
50%        37.000000
75%        46.000000
max        90.000000
Name: age, dtype: float64
Race: Black, sex Male
count    1569.000000
mean       37.682600
std        12.882612
min        17.000000
25%        27.000000
50%        36.000000
75%        46.000000
max        90.000000
Name: age, dtype: float64
Race: Other, sex Female
```

```
count      109.000000
mean        31.678899
std         11.631599
min         17.000000
25%         23.000000
50%         29.000000
75%         39.000000
max         74.000000
Name: age, dtype: float64
Race: Other, sex Male
count      162.000000
mean        34.654321
std         11.355531
min         17.000000
25%         26.000000
50%         32.000000
75%         42.000000
max         77.000000
Name: age, dtype: float64
Race: White, sex Female
count     8642.000000
mean        36.811618
std         14.329093
min         17.000000
25%         25.000000
50%         35.000000
75%         46.000000
max         90.000000
Name: age, dtype: float64
Race: White, sex Male
count    19174.000000
mean        39.652498
std         13.436029
min         17.000000
25%         29.000000
50%         38.000000
75%         49.000000
max         90.000000
Name: age, dtype: float64
```

1. Among whom the proportion of those who earn a lot(>50K) is more: among married or single men (marital-status feature)? Consider married those who have a marital-status starting with Married (Married-civ-spouse, Married-spouse-absent or Married-AF-spouse), the rest are considered bachelors.

In [118]:

```python
data.loc[
    (data['sex'] == 'Male') &
    (data['marital-status'].isin(['Never-married', 'Separated', 'Divorced','Widowed'
])), 'salary'
].value_counts()
```

Out[118]:

```
<=50K    7552
>50K      697
Name: salary, dtype: int64
```

```
data.loc[(data['sex'] == 'Male') &
     (data['marital-status'].str.startswith('Married')), 'salary'].value_counts()
```

Out[36]:

```
<=50K     7576
>50K      5965
Name: salary, dtype: int64
```

In [37]:

```
data['marital-status'].value_counts()
```

Out[37]:

```
Married-civ-spouse       14976
Never-married            10683
Divorced                  4443
Separated                 1025
Widowed                    993
Married-spouse-absent      418
Married-AF-spouse           23
Name: marital-status, dtype: int64
```

1. What is the maximum number of hours a person works per week (hours-per-week feature)? How many people work such a number of hours and what is the percentage of those who earn a lot among them?

In [38]:

```
max_load = data['hours-per-week'].max()
print("Max time - {0} hours./week.".format(max_load))

num_workaholics = data[data['hours-per-week'] == max_load].shape[0]
print("Total number of such hard workers {0}".format(num_workaholics))

rich_share = float(data[(data['hours-per-week'] == max_load)
                & (data['salary'] == '>50K')].shape[0]) / num_workaholics
print("Percentage of rich among them {0}%".format(int(100 * rich_share)))
```

```
Max time - 99 hours./week.
Total number of such hard workers 85
Percentage of rich among them 29%
```

1. Count the average time of work (hours-per-week) those who earning a little and a lot (salary) for each country (native-country).

```python
for (country, salary), sub_df in data.groupby(['native-country', 'salary']):
    print(country, salary, round(sub_df['hours-per-week'].mean(), 2))
```

```
? <=50K 40.16
? >50K 45.55
Cambodia <=50K 41.42
Cambodia >50K 40.0
Canada <=50K 37.91
Canada >50K 45.64
China <=50K 37.38
China >50K 38.9
Columbia <=50K 38.68
Columbia >50K 50.0
Cuba <=50K 37.99
Cuba >50K 42.44
Dominican-Republic <=50K 42.34
Dominican-Republic >50K 47.0
Ecuador <=50K 38.04
Ecuador >50K 48.75
El-Salvador <=50K 36.03
El-Salvador >50K 45.0
England <=50K 40.48
England >50K 44.53
France <=50K 41.06
France >50K 50.75
Germany <=50K 39.14
Germany >50K 44.98
Greece <=50K 41.81
Greece >50K 50.62
Guatemala <=50K 39.36
Guatemala >50K 36.67
Haiti <=50K 36.33
Haiti >50K 42.75
Holand-Netherlands <=50K 40.0
Honduras <=50K 34.33
Honduras >50K 60.0
Hong <=50K 39.14
Hong >50K 45.0
Hungary <=50K 31.3
Hungary >50K 50.0
India <=50K 38.23
India >50K 46.48
Iran <=50K 41.44
Iran >50K 47.5
Ireland <=50K 40.95
Ireland >50K 48.0
Italy <=50K 39.62
Italy >50K 45.4
Jamaica <=50K 38.24
Jamaica >50K 41.1
Japan <=50K 41.0
Japan >50K 47.96
Laos <=50K 40.38
Laos >50K 40.0
Mexico <=50K 40.0
Mexico >50K 46.58
Nicaragua <=50K 36.09
Nicaragua >50K 37.5
Outlying-US(Guam-USVI-etc) <=50K 41.86
Peru <=50K 35.07
Peru >50K 40.0
Philippines <=50K 38.07
Philippines >50K 43.03
Poland <=50K 38.17
```

```
Poland >50K 39.0
Portugal <=50K 41.94
Portugal >50K 41.5
Puerto-Rico <=50K 38.47
Puerto-Rico >50K 39.42
Scotland <=50K 39.44
Scotland >50K 46.67
South <=50K 40.16
South >50K 51.44
Taiwan <=50K 33.77
Taiwan >50K 46.8
Thailand <=50K 42.87
Thailand >50K 58.33
Trinadad&Tobago <=50K 37.06
Trinadad&Tobago >50K 40.0
United-States <=50K 38.8
United-States >50K 45.51
Vietnam <=50K 37.19
Vietnam >50K 39.2
Yugoslavia <=50K 41.6
Yugoslavia >50K 49.5
```

In [42]:

```python
pd.crosstab(data['native-country'], data['salary'],
            values=data['hours-per-week'], aggfunc=np.mean).T
```

Out[42]:

native-country	?	Cambodia	Canada	China	Columbia	Cuba	Dominican-Republic	Ec
salary								
<=50K	40.164760	41.416667	37.914634	37.381818	38.684211	37.985714	42.338235	38.0
>50K	45.547945	40.000000	45.641026	38.900000	50.000000	42.440000	47.000000	48.7

3.2 Часть 2

Использование Pandas для запросов на соединение и группировку данных

In [44]:

```python
user_usage = pd.read_csv("user_usage.csv")
user_device = pd.read_csv("user_device.csv")
devices = pd.read_csv("android_devices.csv")
```

```
user_usage.head()
```

	outgoing_mins_per_month	outgoing_sms_per_month	monthly_mb	use_id
0	21.97	4.82	1557.33	22787
1	1710.08	136.88	7267.55	22788
2	1710.08	136.88	7267.55	22789
3	94.46	35.17	519.12	22790
4	71.59	79.26	1557.33	22792

```
user_device.head(3)
```

	use_id	user_id	platform	platform_version	device	use_type_id
0	22782	26980	ios	10.2	iPhone7,2	2
1	22783	29628	android	6.0	Nexus 5	3
2	22784	28473	android	5.1	SM-G903F	1

```
devices.head(10)
```

	Retail Branding	Marketing Name	Device	Model
0	NaN	NaN	AD681H	Smartfren Andromax AD681H
1	NaN	NaN	FJL21	FJL21
2	NaN	NaN	T31	Panasonic T31
3	NaN	NaN	hws7721g	MediaPad 7 Youth 2
4	3Q	OC1020A	OC1020A	OC1020A
5	7Eleven	IN265	IN265	IN265
6	A.O.I. ELECTRONICS FACTORY	A.O.I.	TR10CS1_11	TR10CS1
7	AG Mobile	AG BOOST 2	BOOST2	E4010
8	AG Mobile	AG Flair	AG_Flair	Flair
9	AG Mobile	AG Go Tab Access 2	AG_Go_Tab_Access_2	AG_Go_Tab_Access_2

```
result = pd.merge(user_usage,
                  user_device[['use_id', 'platform', 'device']],
                  on='use_id')
result.head()
```

	outgoing_mins_per_month	outgoing_sms_per_month	monthly_mb	use_id	platform	devic
0	21.97	4.82	1557.33	22787	android	GT I950
1	1710.08	136.88	7267.55	22788	android	SM G930
2	1710.08	136.88	7267.55	22789	android	SM G930
3	94.46	35.17	519.12	22790	android	D230
4	71.59	79.26	1557.33	22792	android	SM G361

```
print("user_usage dimensions: {}".format(user_usage.shape))
print("user_device dimensions: {}".format(user_device[['use_id', 'platform', 'device']]
.shape))
```

```
user_usage dimensions: (240, 4)
user_device dimensions: (272, 3)
```

```
user_usage['use_id'].isin(user_device['use_id']).value_counts()
```

```
True      159
False      81
Name: use_id, dtype: int64
```

```
result = pd.merge(user_usage,
                  user_device[['use_id', 'platform', 'device']],
                  on='use_id', how='left')
print("user_usage dimensions: {}".format(user_usage.shape))
print("result dimensions: {}".format(result.shape))
print("There are {} missing values in the result.".format(
        result['device'].isnull().sum()))
```

```
user_usage dimensions: (240, 4)
result dimensions: (240, 6)
There are 81 missing values in the result.
```

```
In [53]:
```

```
result.head()
```

```
Out[53]:
```

	outgoing_mins_per_month	outgoing_sms_per_month	monthly_mb	use_id	platform	devic
0	21.97	4.82	1557.33	22787	android	G1 I950
1	1710.08	136.88	7267.55	22788	android	SM G930
2	1710.08	136.88	7267.55	22789	android	SM G930
3	94.46	35.17	519.12	22790	android	D230
4	71.59	79.26	1557.33	22792	android	SM G361

◀ ▬▬▬▬▬▬▬▬▬▬▬▬▬▬▬ ▶

```
In [54]:
```

```
result.tail()
```

```
Out[54]:
```

	outgoing_mins_per_month	outgoing_sms_per_month	monthly_mb	use_id	platform	dev
235	260.66	68.44	896.96	25008	NaN	
236	97.12	36.50	2815.00	25040	NaN	
237	355.93	12.37	6828.09	25046	NaN	
238	632.06	120.46	1453.16	25058	NaN	
239	488.70	906.92	3089.85	25220	NaN	

◀ ▬▬▬▬▬▬▬▬▬▬▬▬▬▬▬ ▶

```
In [55]:
```

```
result = pd.merge(user_usage,
                  user_device[['use_id', 'platform', 'device']],
                  on='use_id', how='right')
print("user_device dimensions: {}".format(user_device.shape))
print("result dimensions: {}".format(result.shape))
print("There are {} missing values in the 'monthly_mb' column in the result.".format(
        result['monthly_mb'].isnull().sum()))
print("There are {} missing values in the 'platform' column in the result.".format(
        result['platform'].isnull().sum()))
```

```
user_device dimensions: (272, 6)
result dimensions: (272, 6)
There are 113 missing values in the 'monthly_mb' column in the result.
There are 0 missing values in the 'platform' column in the result.
```

In [56]:

```python
print("There are {} unique values of use_id in our dataframes.".format(
        pd.concat([user_usage['use_id'], user_device['use_id']]).unique().shape[0]))
result = pd.merge(user_usage,
                  user_device[['use_id', 'platform', 'device']],
                  on='use_id', how='outer', indicator=True)

print("Outer merge result has {} rows.".format(result.shape))

print("There are {} rows with no missing values.".format(
    (result.apply(lambda x: x.isnull().sum(), axis=1) == 0).sum()))
```

There are 353 unique values of use_id in our dataframes.
Outer merge result has (353, 7) rows.
There are 159 rows with no missing values.

In [57]:

```python
result.iloc[[0, 1, 200,201, 350,351]]
```

Out[57]:

	outgoing_mins_per_month	outgoing_sms_per_month	monthly_mb	use_id	platform	
0	21.97	4.82	1557.33	22787	android	G⁻
1	1710.08	136.88	7267.55	22788	android	
200	28.79	29.42	3114.67	23988	NaN	
201	616.56	99.85	5414.14	24006	NaN	
350	NaN	NaN	NaN	23050	ios	iPh
351	NaN	NaN	NaN	23051	ios	iPh

In [58]:

```python
# First, add the platform and device to the user usage.
result = pd.merge(user_usage,
                  user_device[['use_id', 'platform', 'device']],
                  on='use_id',
                  how='left')

# Now, based on the "device" column in result, match the "Model" column in devices.
devices.rename(columns={"Retail Branding": "manufacturer"}, inplace=True)
result = pd.merge(result,
                  devices[['manufacturer', 'Model']],
                  left_on='device',
                  right_on='Model',
                  how='left')

result.head()
```

Out[58]:

	outgoing_mins_per_month	outgoing_sms_per_month	monthly_mb	use_id	platform	devic
0	21.97	4.82	1557.33	22787	android	GT I950
1	1710.08	136.88	7267.55	22788	android	SM G930
2	1710.08	136.88	7267.55	22789	android	SM G930
3	94.46	35.17	519.12	22790	android	D230
4	71.59	79.26	1557.33	22792	android	SM G361

In [59]:

```python
devices[devices.Model == 'SM-G930F']
```

Out[59]:

	manufacturer	Marketing Name	Device	Model
10381	Samsung	Galaxy S7	herolte	SM-G930F

```python
devices[devices.Device.str.startswith('GT')]
```

	manufacturer	Marketing Name	Device	Model
1095	Bitmore	GTAB700	GTAB700	NID_7010
1096	Bitmore	GTAB900	GTAB900	S952
2402	Grundig	GTB1050	GTB1050	GTB 1050
2403	Grundig	GTB850	GTB850	GTB 850
2404	Grundig	TC69CA2	GTB801	GTB 801
9125	Samsung	NaN	GT-I5510M	GT-I5510M
9126	Samsung	NaN	GT-I5510T	GT-I5510T
9127	Samsung	NaN	GT-I5800L	GT-I5800L
9128	Samsung	NaN	GT-N7000B	GT-N7000B
9129	Samsung	NaN	GT-P7300B	GT-P7300B
9130	Samsung	NaN	GT-P7320T	GT-P7320T
9131	Samsung	NaN	GT-P7500M	GT-P7500M
9132	Samsung	NaN	GT-P7500R	GT-P7500R
9133	Samsung	NaN	GT-P7500V	GT-P7500V
9134	Samsung	NaN	GT-S5698	GT-S5698
9135	Samsung	NaN	GT-S5820	GT-S5820
9136	Samsung	NaN	GT-S5830V	GT-S5830V
9173	Samsung	Absolute	GT-B9120	GT-B9120
9184	Samsung	Europa	GT-I5500B	GT-I5500B
9185	Samsung	Europa	GT-I5500L	GT-I5500L
9186	Samsung	Europa	GT-I5500M	GT-I5500M
9187	Samsung	Europa	GT-I5503T	GT-I5503T
9188	Samsung	Europa	GT-I5510L	GT-I5510L
9191	Samsung	Galaxy (China)	GT-B9062	GT-B9062
9288	Samsung	Galaxy Ace	GT-S5830	GT-S5830
9289	Samsung	Galaxy Ace	GT-S5830B	GT-S5830B
9290	Samsung	Galaxy Ace	GT-S5830C	GT-S5830C
9291	Samsung	Galaxy Ace	GT-S5830D	GT-S5830D
9292	Samsung	Galaxy Ace	GT-S5830F	GT-S5830F
9293	Samsung	Galaxy Ace	GT-S5830G	GT-S5830G
...	...	...	...	...
10480	Samsung	Galaxy Tab 7.7	GT-P6800	GT-P6800
10481	Samsung	Galaxy Tab 7.7	GT-P6810	GT-P6810
10484	Samsung	Galaxy Tab 8.9	GT-P7300	GT-P7300
10485	Samsung	Galaxy Tab 8.9	GT-P7310	GT-P7310
10486	Samsung	Galaxy Tab 8.9	GT-P7320	NT-P7320
10778	Samsung	Galaxy W	GT-I8150	GT-I8150

	manufacturer	Marketing Name	Device	Model
10779	Samsung	Galaxy W	GT-I8150B	GT-I8150B
10780	Samsung	Galaxy W	GT-I8150T	GT-I8150T
10796	Samsung	Galaxy Xcover	GT-S5690	GT-S5690
10797	Samsung	Galaxy Xcover	GT-S5690L	GT-S5690L
10798	Samsung	Galaxy Xcover	GT-S5690M	GT-S5690M
10799	Samsung	Galaxy Xcover	GT-S5690R	GT-S5690R
10804	Samsung	Galaxy Y	GT-S5360	GT-S5360
10805	Samsung	Galaxy Y	GT-S5360B	GT-S5360B
10806	Samsung	Galaxy Y	GT-S5360L	GT-S5360L
10807	Samsung	Galaxy Y	GT-S5360T	GT-S5360T
10808	Samsung	Galaxy Y	GT-S5363	GT-S5363
10809	Samsung	Galaxy Y	GT-S5368	GT-S5368
10810	Samsung	Galaxy Y	GT-S5369	GT-S5369
10813	Samsung	Galaxy Y Duos	GT-S6102	GT-S6102
10814	Samsung	Galaxy Y Duos	GT-S6102B	GT-S6102B
10815	Samsung	Galaxy Y Duos	GT-S6102E	GT-S6102E
10818	Samsung	Galaxy Y Pop	GT-S6108	GT-S6108
10819	Samsung	Galaxy Y Pro	GT-B5510	GT-B5510
10820	Samsung	Galaxy Y Pro	GT-B5510B	GT-B5510B
10821	Samsung	Galaxy Y Pro	GT-B5510L	GT-B5510L
10822	Samsung	Galaxy Y Pro Duos	GT-B5512	GT-B5512
10823	Samsung	Galaxy Y Pro Duos	GT-B5512B	GT-B5512B
10824	Samsung	Galaxy Y TV	GT-S5367	GT-S5367
10979	Sharp	AQUOS SERIE mini SHV38	GTQ	SHV38

164 rows × 4 columns

```
result.head()
```

Out[61]:

	outgoing_mins_per_month	outgoing_sms_per_month	monthly_mb	use_id	platform	devic
0	21.97	4.82	1557.33	22787	android	GT I950
1	1710.08	136.88	7267.55	22788	android	SM G930
2	1710.08	136.88	7267.55	22789	android	SM G930
3	94.46	35.17	519.12	22790	android	D230
4	71.59	79.26	1557.33	22792	android	SM G361

In [62]:

```
result.groupby("manufacturer").agg({
        "outgoing_mins_per_month": "mean",
        "outgoing_sms_per_month": "mean",
        "monthly_mb": "mean",
        "use_id": "count"
    })
```

Out[62]:

manufacturer	outgoing_mins_per_month	outgoing_sms_per_month	monthly_mb	use_id
HTC	299.842955	93.059318	5144.077955	44
Huawei	81.526667	9.500000	1561.226667	3
LGE	111.530000	12.760000	1557.330000	2
Lava	60.650000	261.900000	12458.670000	2
Lenovo	215.920000	12.930000	1557.330000	2
Motorola	95.127500	65.666250	3946.500000	16
OnePlus	354.855000	48.330000	6575.410000	6
Samsung	191.010093	92.390463	4017.318889	108
Sony	177.315625	40.176250	3212.000625	16
Vodafone	42.750000	46.830000	5191.120000	1
ZTE	42.750000	46.830000	5191.120000	1

Использование Pandasql для запросов на соединение и группировку данных

```
import pandasql as ps
```

In [110]:

```
# pandasql code
def left_join_ps(user_usage, user_device):
    join_query = '''
    SELECT user_usage.outgoing_mins_per_month,
    user_usage.outgoing_sms_per_month,
    user_usage.monthly_mb,
    user_device.use_id,
    user_device.device,
    user_device.platform

    FROM user_usage
    LEFT JOIN user_device
    ON user_usage.use_id = user_device.use_id;
    '''
    return ps.sqldf(join_query, locals())
left_join_ps(user_usage, user_device).tail()
```

Out[110]:

	outgoing_mins_per_month	outgoing_sms_per_month	monthly_mb	use_id	device	platf
235	260.66	68.44	896.96	NaN	None	N
236	97.12	36.50	2815.00	NaN	None	N
237	355.93	12.37	6828.09	NaN	None	N
238	632.06	120.46	1453.16	NaN	None	N
239	488.70	906.92	3089.85	NaN	None	N

Использование функций агрегирования

In [147]:

```
# pandasql code
def aggregating_ps(user_usage, user_device):
    aggr_query = '''
        SELECT
            avg(outgoing_mins_per_month) as outgoing_mins_per_month,
            user_device.platform
        FROM user_usage
        LEFT JOIN user_device
        ON user_usage.use_id = user_device.use_id
        GROUP BY platform
        '''
    return ps.sqldf(aggr_query, locals())
```

```
aggregating_ps(user_usage, user_device)
```

	outgoing_mins_per_month	platform
0	414.376420	None
1	201.258535	android
2	366.060000	ios

```
merged = pd.merge(user_usage,
                  user_device[['use_id', 'platform']],
                  on='use_id',
                  how='left')
merged['platform'] = merged['platform'].astype(str)
merged.groupby('platform').mean()
```

	outgoing_mins_per_month	outgoing_sms_per_month	monthly_mb	use_id
platform				
android	201.258535	85.354586	4221.387834	22922.350318
ios	366.060000	293.975000	961.155000	22920.500000
nan	414.376420	120.540370	2545.485062	23998.444444

Сравнение времени выполнения каждого запроса в Pandas и PandaSQL.

```python
import time

def count_mean_time(func, params, N =5):
    total_time = 0
    for i in range(N):
        time1 = time.time()
        if len(params) == 1:
            tmp_df = func(params[0])
        elif len(params) == 2:
            tmp_df = func(params[0], params[1])
        time2 = time.time()
        total_time += (time2 - time1)
    return total_time/N
```

```
left_join_pandasql_time = count_mean_time(left_join_ps, [user_usage, user_device], N=20
)
left_join_pandasql_time
```

Out[113]:

0.019597113132476807

```
def left_join_pandas(user_usage, user_device):
    result = pd.merge(user_usage,
                user_device[['use_id', 'platform', 'device']],
                on='use_id', how='left')
    return result
left_join_pandas_time = count_mean_time(left_join_pandas, [user_usage, user_device], N=
20)
left_join_pandas_time
```

Out[114]:

0.0056846380233764645

```
aggreg_pandasql_time = count_mean_time(aggregating_ps, [user_usage, user_device], N=20)
aggreg_pandasql_time
```

Out[115]:

0.017153775691986083

```
def aggreg_query_pandas(user_usage, user_device):
    merged = pd.merge(user_usage,
                user_device[['use_id', 'platform']],
                on='use_id',
                how='left')
    return merged.groupby('platform').mean()

aggr_pandas_time = count_mean_time(aggreg_query_pandas, [user_usage, user_device], N=20
)
aggr_pandas_time
```

Out[116]:

0.00812966823577881

```
print('Разница во времени при выполнении запроса LEFT JOIN: {0}'.format(left_join_panda
sql_time - left_join_pandas_time))
print('Разница во времени при выполнении запроса с ф-ями агрегирования: {0}'.format(agg
reg_pandasql_time - aggr_pandas_time))
```

Разница во времени при выполнении запроса LEFT JOIN: 0.013912475109100342
Разница во времени при выполнении запроса с ф-ями агрегирования: 0.0090241
07456207273

Сравнив время выполнения запросов при помощи библиотек Pandas и Pandasql, можно сделать вывод, что для простых запросов лучше использовать Pandas, так как время выполнения запросов к источнику данных меньше, чем у Pandasql. Но при написании сложных запросов Pandasql даст возможность использовать привычные SQL запросы.

Список литературы

1. Гапанюк Ю. Е. Лабораторная работа «Изучение библиотек обработки данных» [Электронный ресурс] // GitHub. – 2019. – Режим доступа: https://github.com/ (https://github.com/) ugapanyuk/ml_course/wiki/LAB_PANDAS (дата обращения: 20.02.2019).
2. pandas 0.24.1 documentation [Electronic resource] // PyData. — 2019. — Access mode: http://pandas.pydata.org/pandas-docs/stable/ (http://pandas.pydata.org/pandas-docs/stable/) (online; accessed: 20.02.2019).
3. You are my Sunshine [Electronic resource] // Space Apps Challenge. — 2017. — Access mode: https://2017.spaceappschallenge.org/challenges/earth-and-us/ (https://2017.spaceappschallenge.org/challenges/earth-and-us/) you-are-my-sunshine/details (online; accessed: 22.02.2019).
4. yhat/pandasql: sqldf for pandas [Electronic resource] // GitHub. — 2017. — Access mode: https://github.com/yhat/pandasql (https://github.com/yhat/pandasql) (online; accessed: 22.02.2019).
5. Team The IPython Development. IPython 7.3.0 Documentation [Electronic resource] // Read the Docs. — 2019. — Access mode: https://ipython.readthedocs.io/en/ (https://ipython.readthedocs.io/en/) stable/ (online; accessed: 20.02.2019).